

Practical Memory Safety with *REST*

KANAD SINHA & SIMHA SETHUMADHAVAN

COLUMBIA UNIVERSITY

A solid orange horizontal bar at the bottom of the slide.

Is memory safety relevant?

Millions of IoT devices hit by 'Devil's Ivy' bug in open source code

Devil's Ivy is likely to remain unpatched

By Liam Tung | July 20, 2017 -- 10:33 GMT (03:00)

Symantec Antivirus products vulnerable to horrid overflow bug

A vulnerability in Symantec's Antivirus products allows for kernel memory corruption without user action on the part of the user.

Zero-day Skype remote code execution

Heartbleed bug still affects thousands of sites

In 2017, 55% of remote-code execution causing bugs in Microsoft due to memory errors

'90s-style security flaws at risk

gear, D-Link, TP-Link devices, and more

Extremely large number of software are vulnerable

2008, vulnerability has left apps and hardware open to remote hijacking.

Bigger than Heartbleed, 'Venom' security vulnerability threatens most datacenters

Security researchers say the zero-day flaw affects 'millions' of machines in datacenters around the world.



By Zack Whittaker for Zero Day | May 13, 2015 -- 12:00 GMT (05:00 PDT) | Topic: Cloud

The bugs could...

data and compromise patient care.

Is memory safety relevant?

Yes!

Practical memory safety

Presenting...

Random Embedded Security Tokens or *REST*

Core H/W primitive: Insert known **64B random value** (*token*) in program and detect accesses to them.

Practical memory safety


Presenting...

Random Embedded Security Tokens or *REST*

Core H/W primitive: Insert known **64B random value (token)** in program and detect accesses to them.

```
char *buf = malloc(BUF_LEN);
```

```
for (i=0; i<out_of_bounds; i++)  
    buf = 0;
```



Heap

Practical memory safety


Presenting...

Random Embedded Security Tokens or *REST*

Core H/W primitive: Insert known **64B random value** (*token*) in program and detect accesses to them.

```
char *buf = malloc(BUF_LEN);
```

```
for (i=0; i<out_of_bounds; i++)  
    buf = 0;
```



Practical memory safety

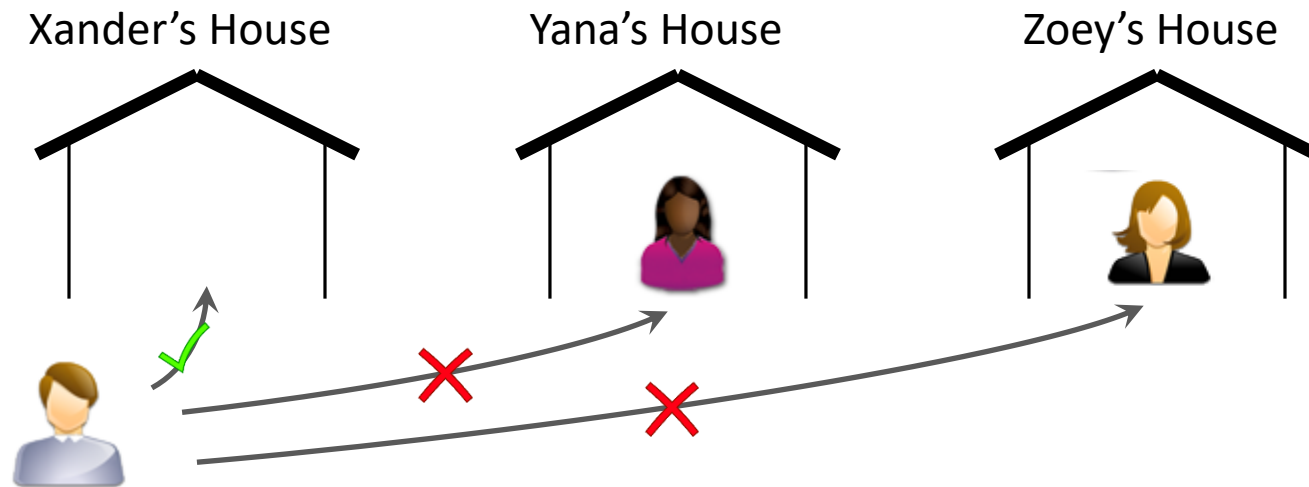
Presenting...

Random Embedded Security Tokens or *REST*

Core H/W primitive: Insert known **64B random value** (*token*) in program and detect accesses to them.

- Trivial hardware implementation
- Software framework based on *AddressSanitizer*
- Provides heap safety for legacy binaries

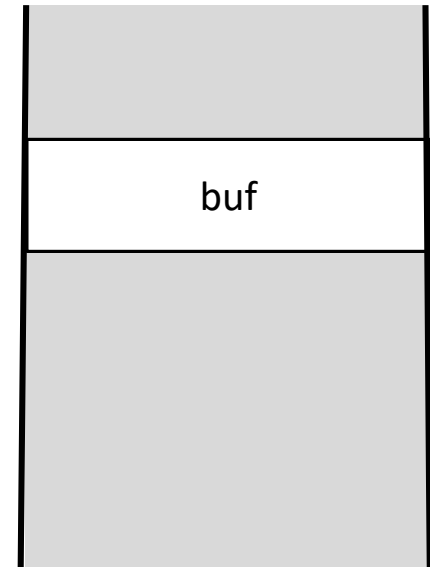
Background: **Spatial** Memory Safety



Background: Spatial Memory Safety

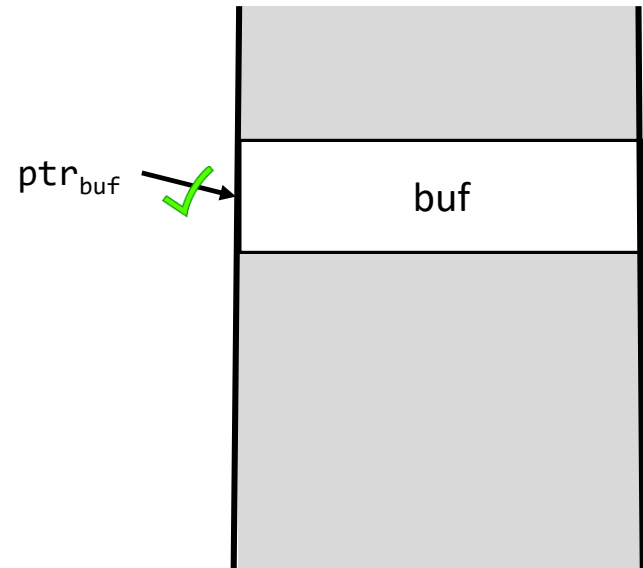
```
char *ptrbuf = malloc(BUF_LEN);  
...  
ptrbuf[in_bounds] = X;  
...  
ptrbuf[out_of_bounds] = Y;
```

ptr_{buf}



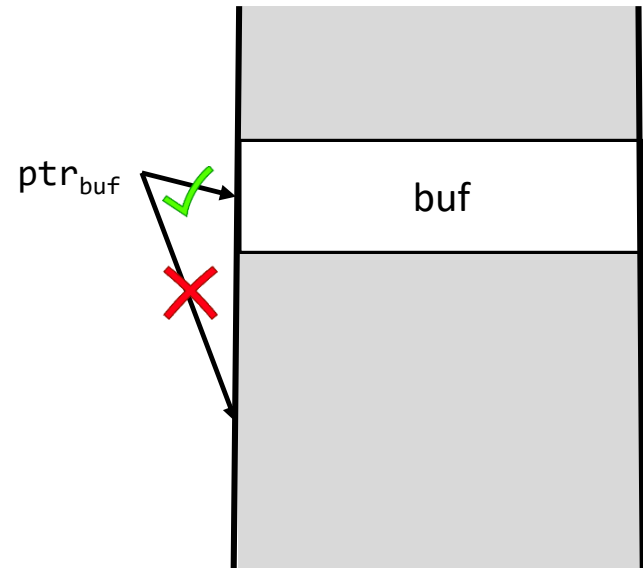
Background: Spatial Memory Safety

```
char *ptr_buf = malloc(BUF_LEN);  
...  
ptr_buf[in_bounds] = X;  
...  
ptr_buf[out_of_bounds] = Y;
```



Background: Spatial Memory Safety

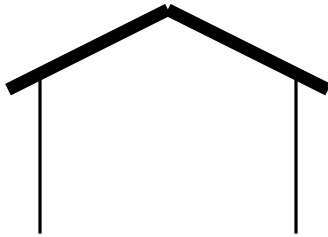
```
char *ptr_buf = malloc(BUF_LEN);  
...  
ptr_buf[in_bounds] = X;  
...  
ptr_buf[out_of_bounds] = Y;
```



Background: Temporal Memory Safety

Xander moves out, Will moves in

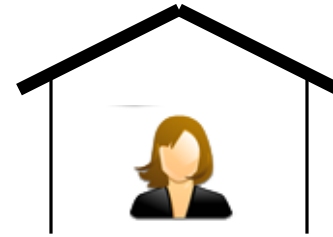
Xander's House



Yana's House

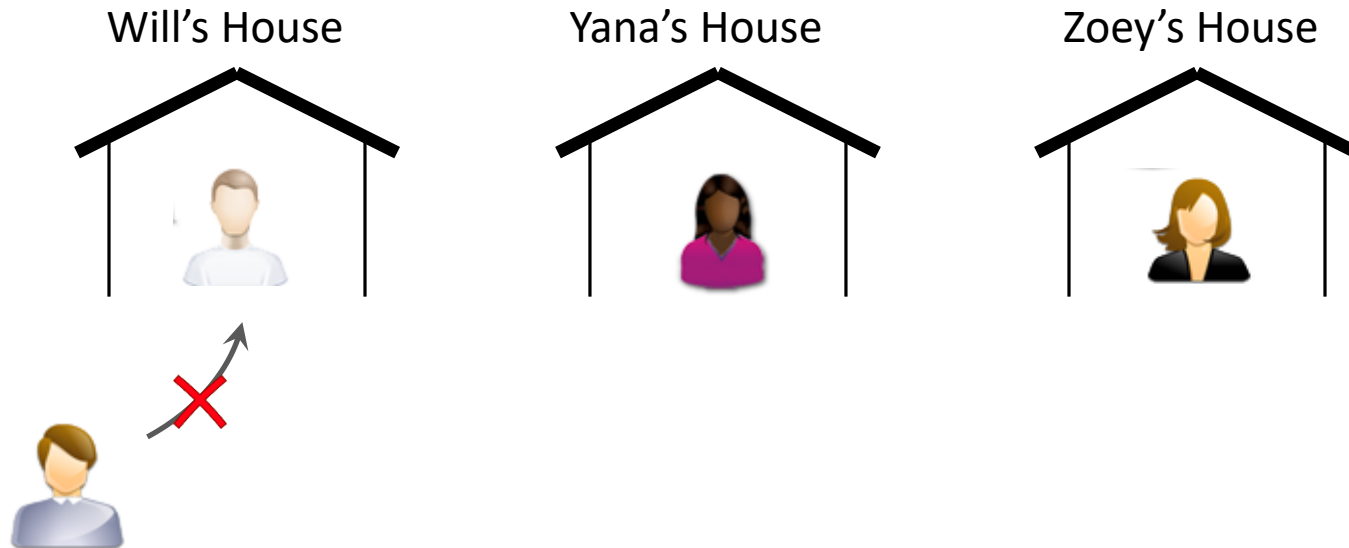


Zoey's House



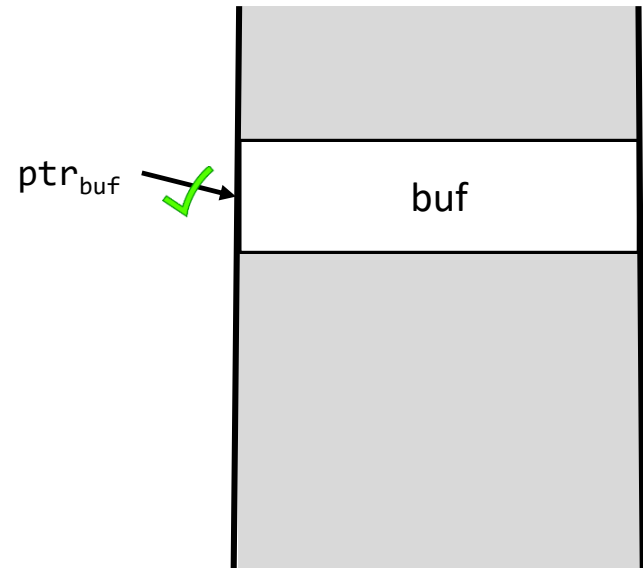
Background: Temporal Memory Safety

Xander moves out, Will moves in



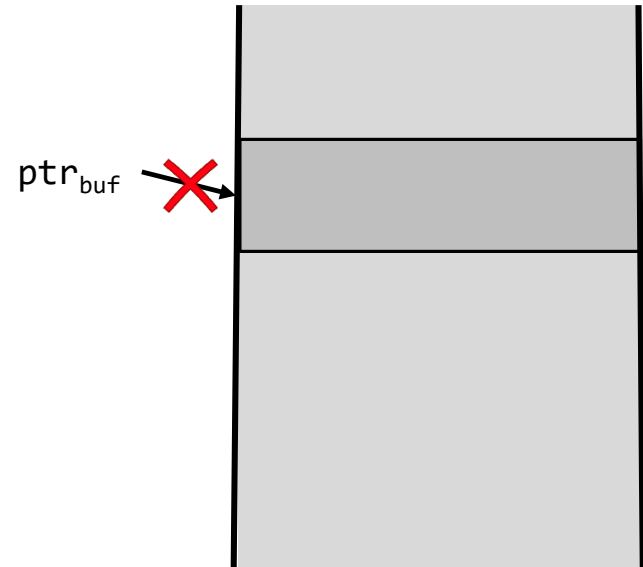
Background: Temporal Memory Safety

```
char *ptrbuf = malloc(BUF_LEN);  
ptrbuf[in_bounds] = X;  
...  
free(ptrbuf);  
ptrbuf[in_bounds] = Y;
```



Background: Temporal Memory Safety

```
char *ptr_buf = malloc(BUF_LEN);  
ptr_buf[in_bounds] = X;  
...  
free(ptr_buf);  
ptr_buf[in_bounds] = Y;
```



Previous H/W Solutions

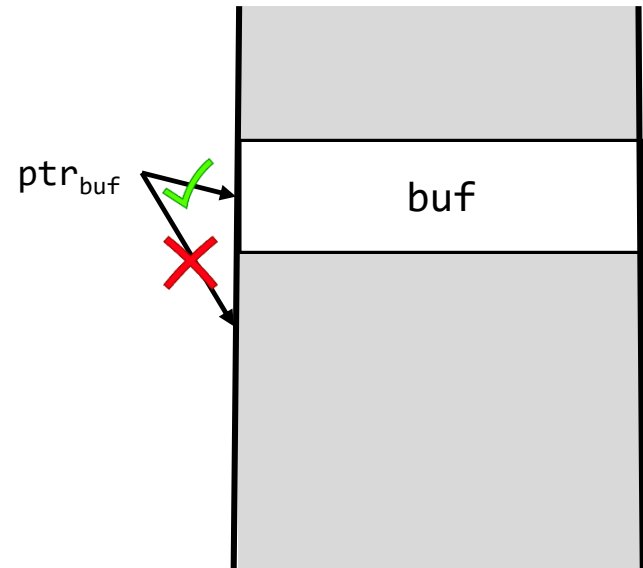
Mainly categorizable into 2 types.

Previous H/W Solutions

Mainly categorizable into 2 types.

- **Whitelisting:** Pointer based


- + Good coverage
- + Temporal safety (for some)
- Performance overhead
- Implementation overhead
- Imprecise



Previous H/W Solutions

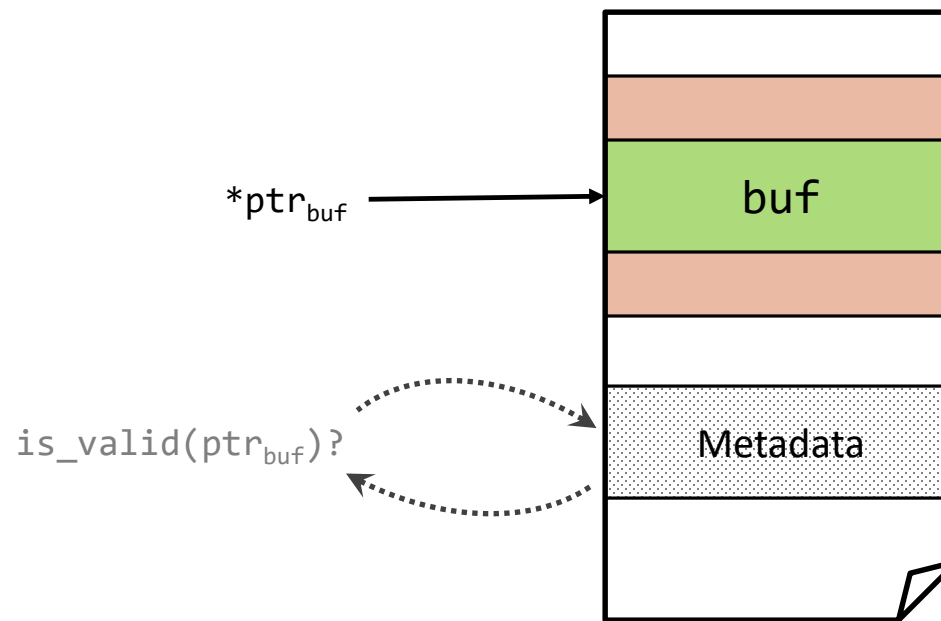
Mainly categorizable into 2 types.

- **Whitelisting:** Pointer based
 - + Good coverage
 - + Temporal safety (for some)
 - Performance overhead
 - Implementation overhead
 - Imprecise
- **Blacklisting:** Location based
 - + Fast
 - Weaker coverage (has false negatives)
 - Implementation overhead
 - No temporal protection



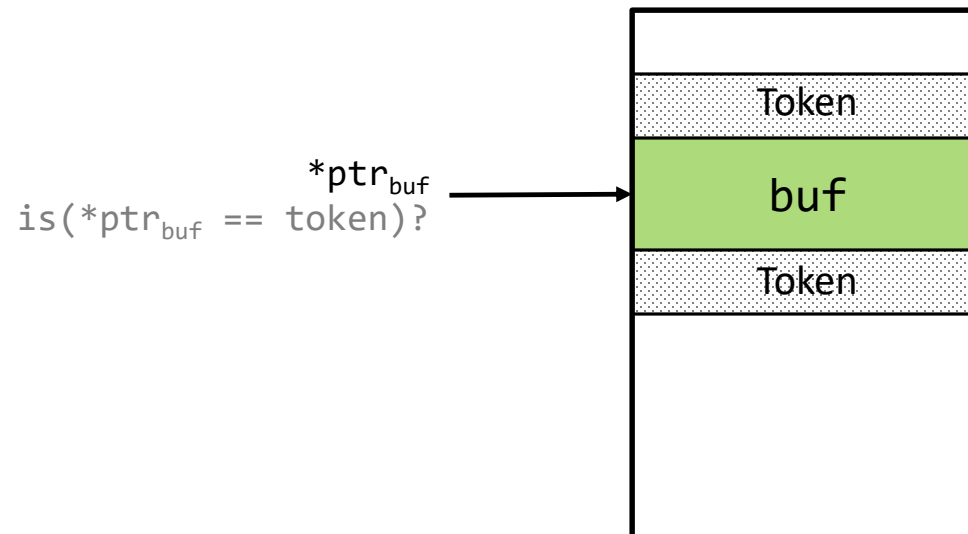
Previous H/W Solutions

Tag-based



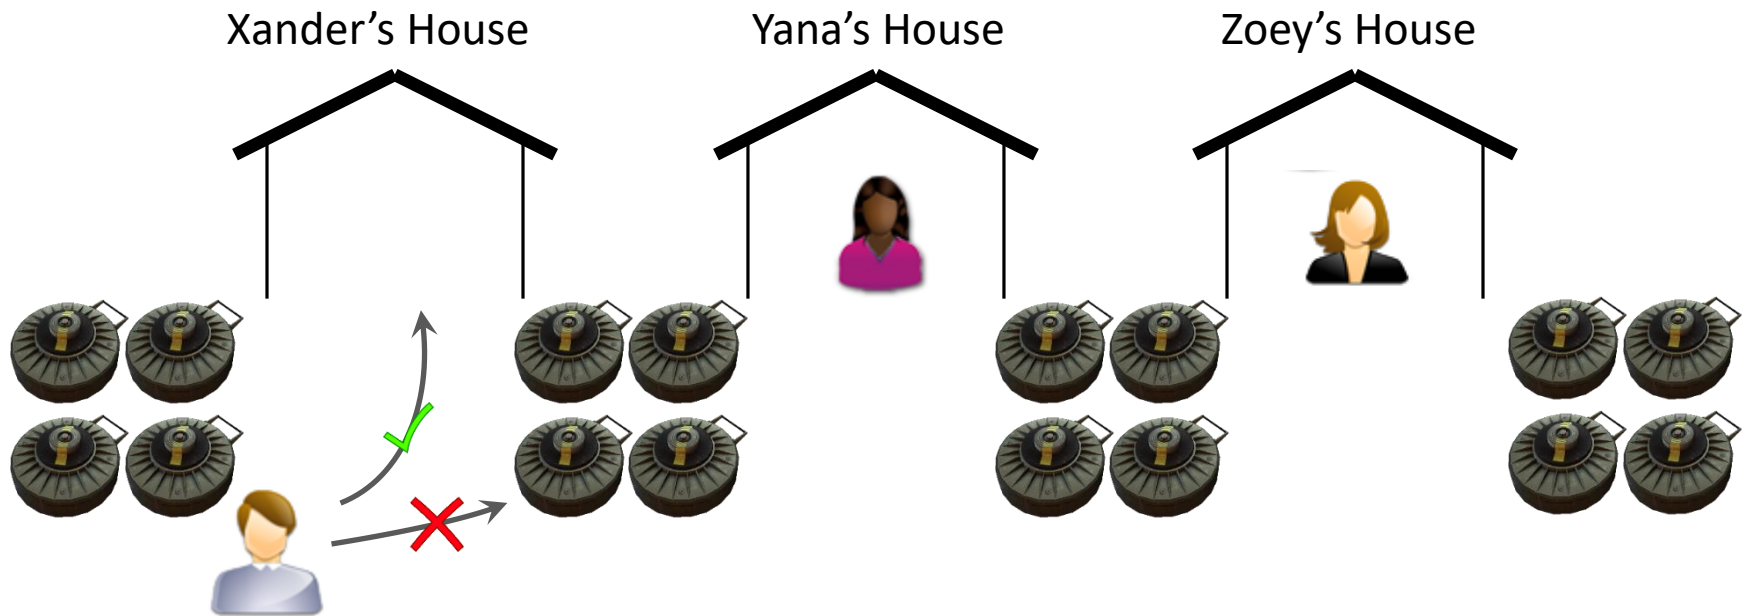
REST: Primitive Overview

Content-based blacklisting



REST primitive has **trivial** complexity, overhead

REST: Spatial Memory Safety



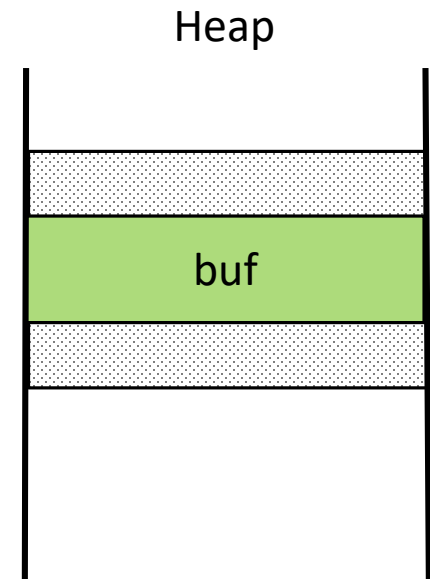
REST: Temporal Memory Safety



REST Software

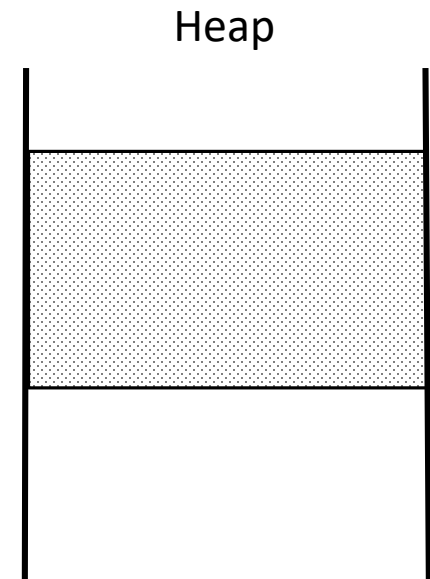
Heap Safety

- Allocate and bookend region, malloc to program



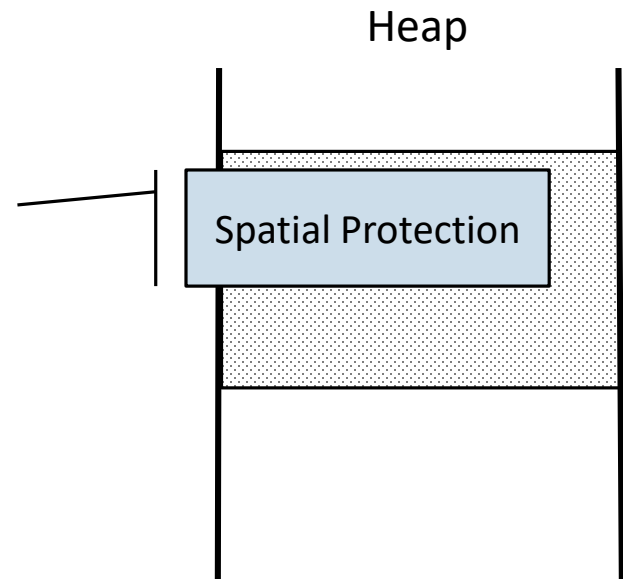
Heap Safety

- Allocate and bookend region, malloc to program
- **REST**'ize at free
- Do not reallocate region until heap sufficiently consumed



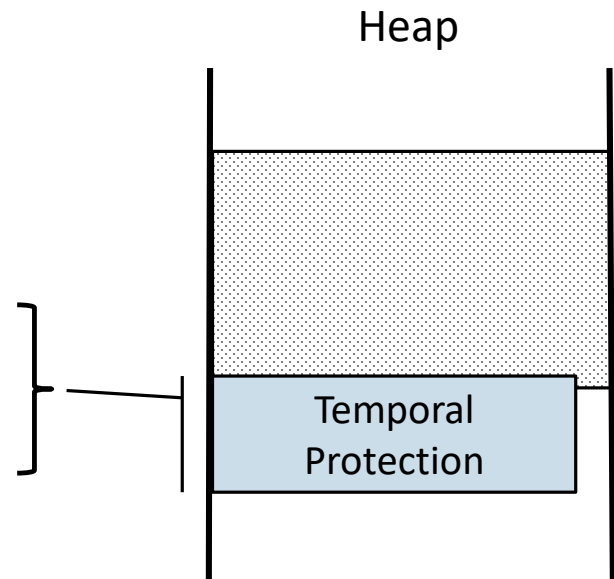
Heap Safety

- Allocate and bookend region, malloc to program
- **REST**'ize at free
- Do not reallocate region until heap sufficiently consumed



Heap Safety


- Allocate and bookend region, malloc to program
- REST'ize at free
- Do not reallocate region until heap sufficiently consumed



Can be enabled for legacy binaries

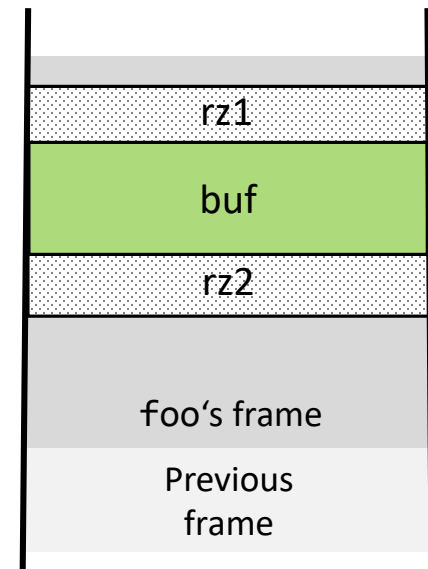
Stack Safety

```
void foo() {  
    char buf[64];  
    ...  
    return;  
}
```




Stack Safety

```
void foo() {  
    char rz1[64];  
    char buf[64];  
    char rz2[64];  
    arm(rz1);  
    arm(rz2);  
    ...  
    disarm(rz1);  
    disarm(rz2);  
    return;  
}
```



Stack Safety

```
void foo() {  
    char rz1[64];  
    char buf[64];  
    char rz2[64];  
    arm(rz1);  
    arm(rz2);  
    ...  
    disarm(rz1);  
    disarm(rz2);  
    return;  
}
```



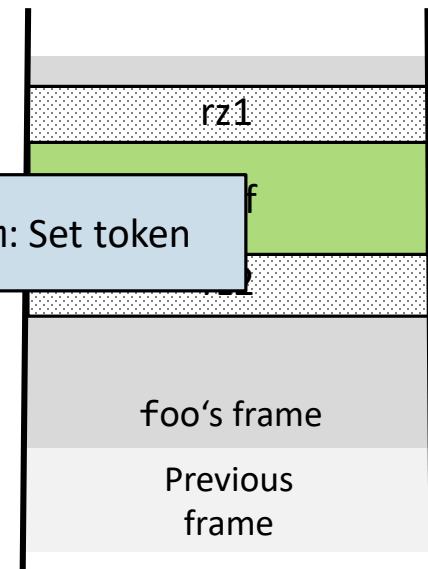
Stack Safety

```
void foo() {  
    char rz1[64];  
    char buf[64];  
    char rz2[64];  
    arm rz1;  
    arm rz2;  
    disarm(rz1);  
    disarm(rz2);  
    return;  
}
```

arm rz1;
arm rz2;

disarm(rz1);
disarm(rz2);
return;

arm: Set token




Stack Safety

```
void foo() {  
    char rz1[64];  
    char buf[64];  
    char rz2[64];  
    arm(rz1);  
    arm(rz2);  
    disarm rz1;  
    disarm rz2;  
    return;  
}
```


disarm rz1;
disarm rz2;

return;



Stack Safety

```
void foo() {  
    char rz1[64];  
    char buf[64];  
    char rz2[64];  
    arm(rz1);  
    arm(rz2);  
    ...  
    disarm(rz1);  
    disarm(rz2);  
    return;  
}
```

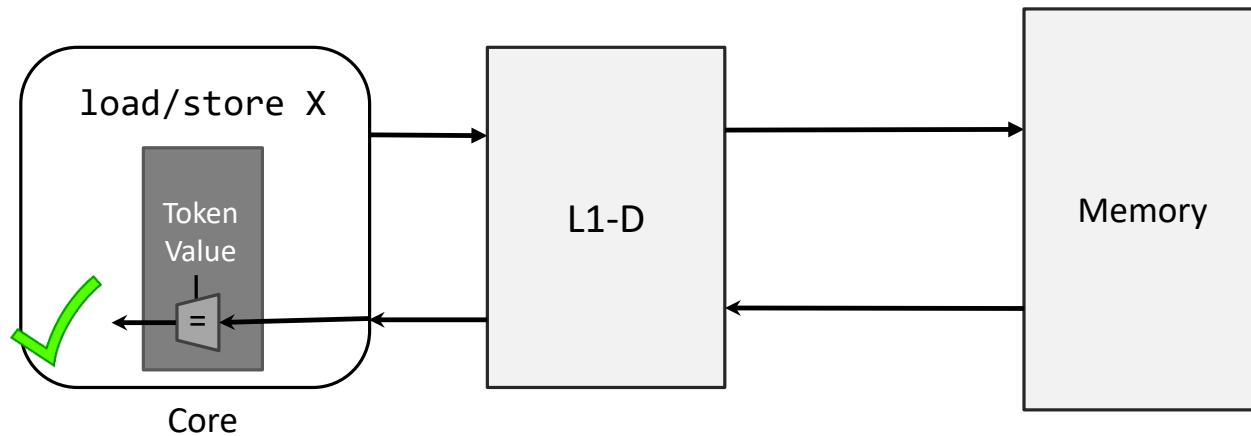


Requires **recompilation** with *REST* plugin

REST Hardware


Naïve Design

Every store involves an extra load → Complicated and expensive




Cache Modifications

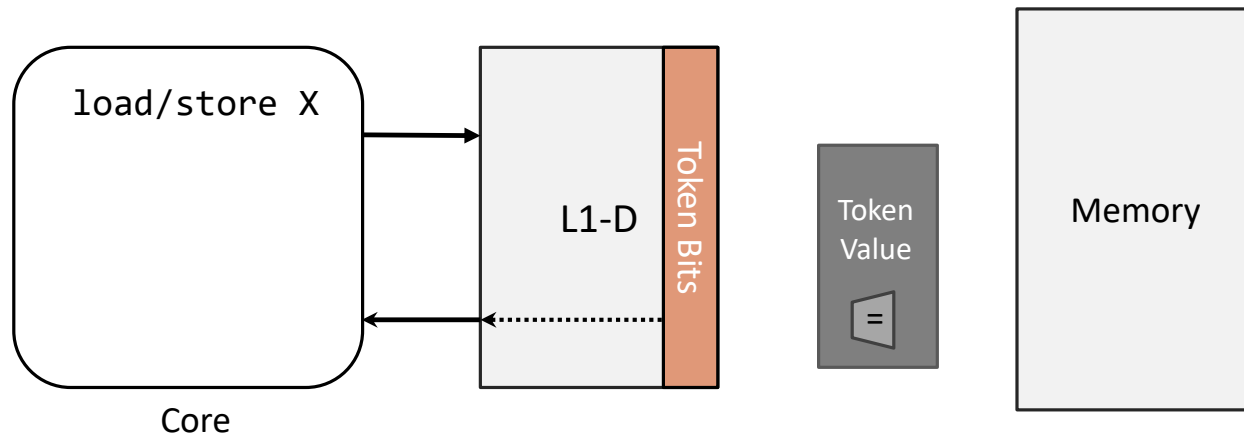
Comparator at L1-D mem interface + **1b** per L1-D line



Cache Miss




Cache Hit



Cache Eviction

Armed outgoing line filled with token value



What about the core?

TODO: Have to support **arms** and **disarms**

- 512b writes
- Special semantics: can only touch token with disarm

LSQ design concerns:

- Forwarding would break semantics
- 512b data entries
- How to match unaligned token access?

Load-Store Queue

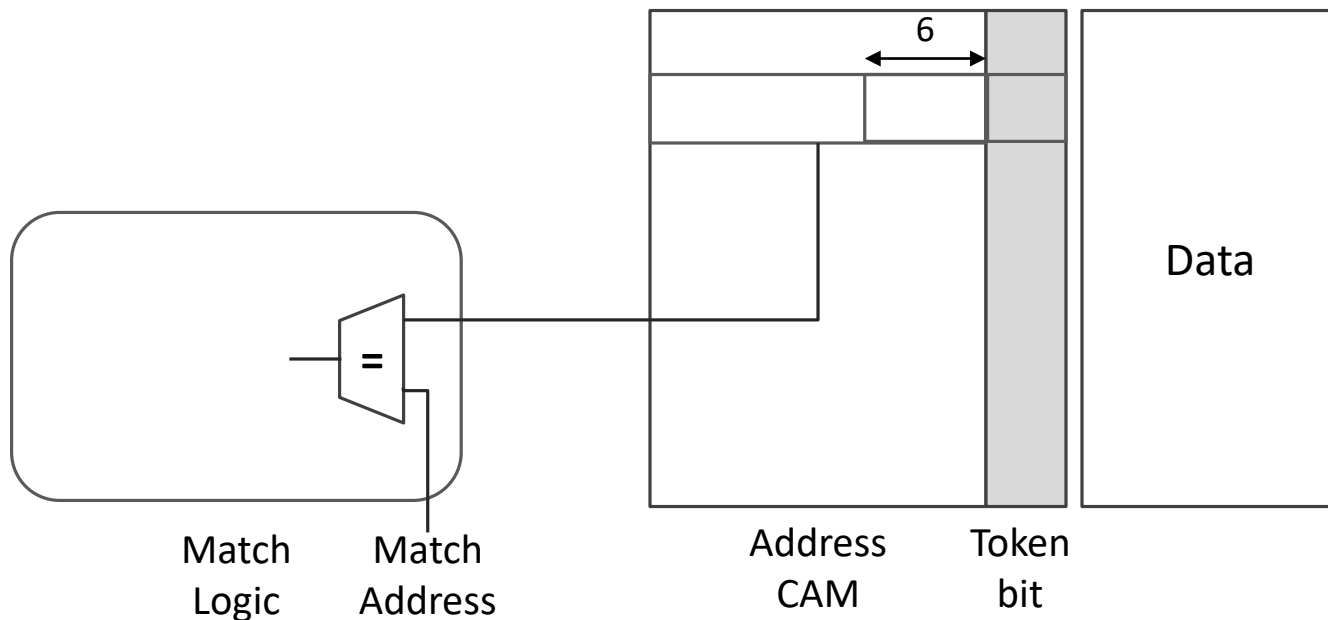
- Forwarding breaks semantics
- 512b data entries
- Detecting unaligned token access

Add 1b tag



Only update token bit

Split regular match logic



Load-Store Queue

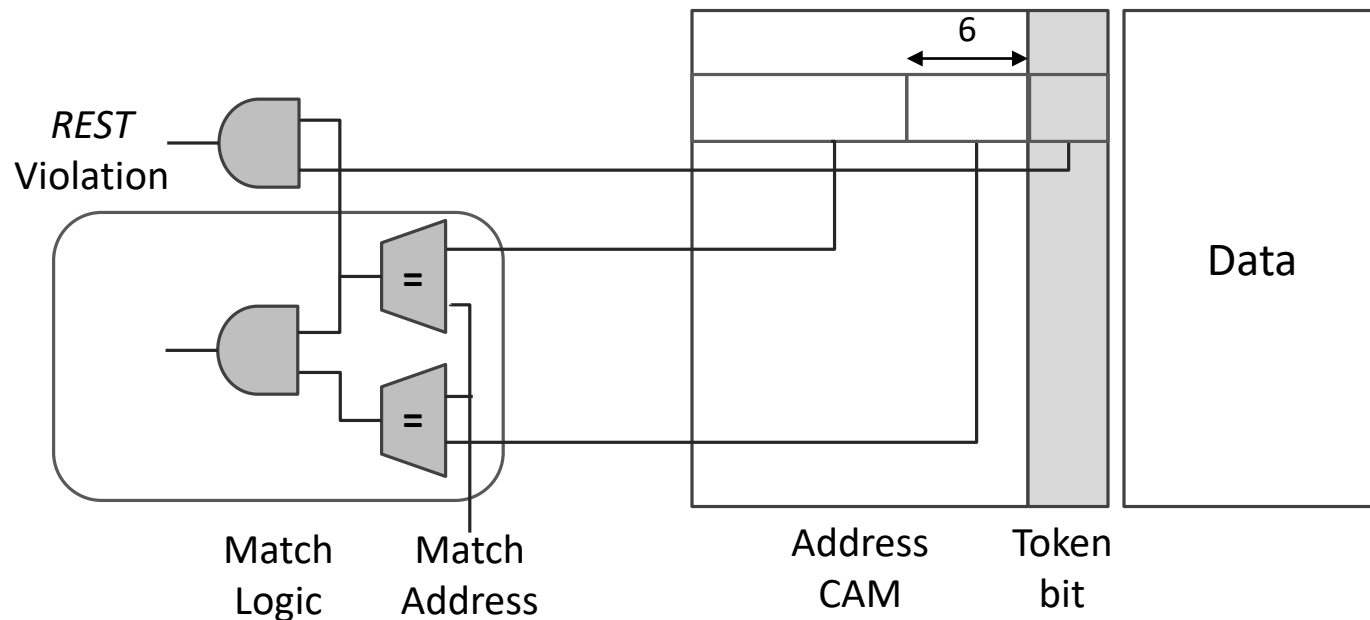
- Forwarding breaks semantics
- 512b data entries
- Detecting unaligned token access

Add 1b tag



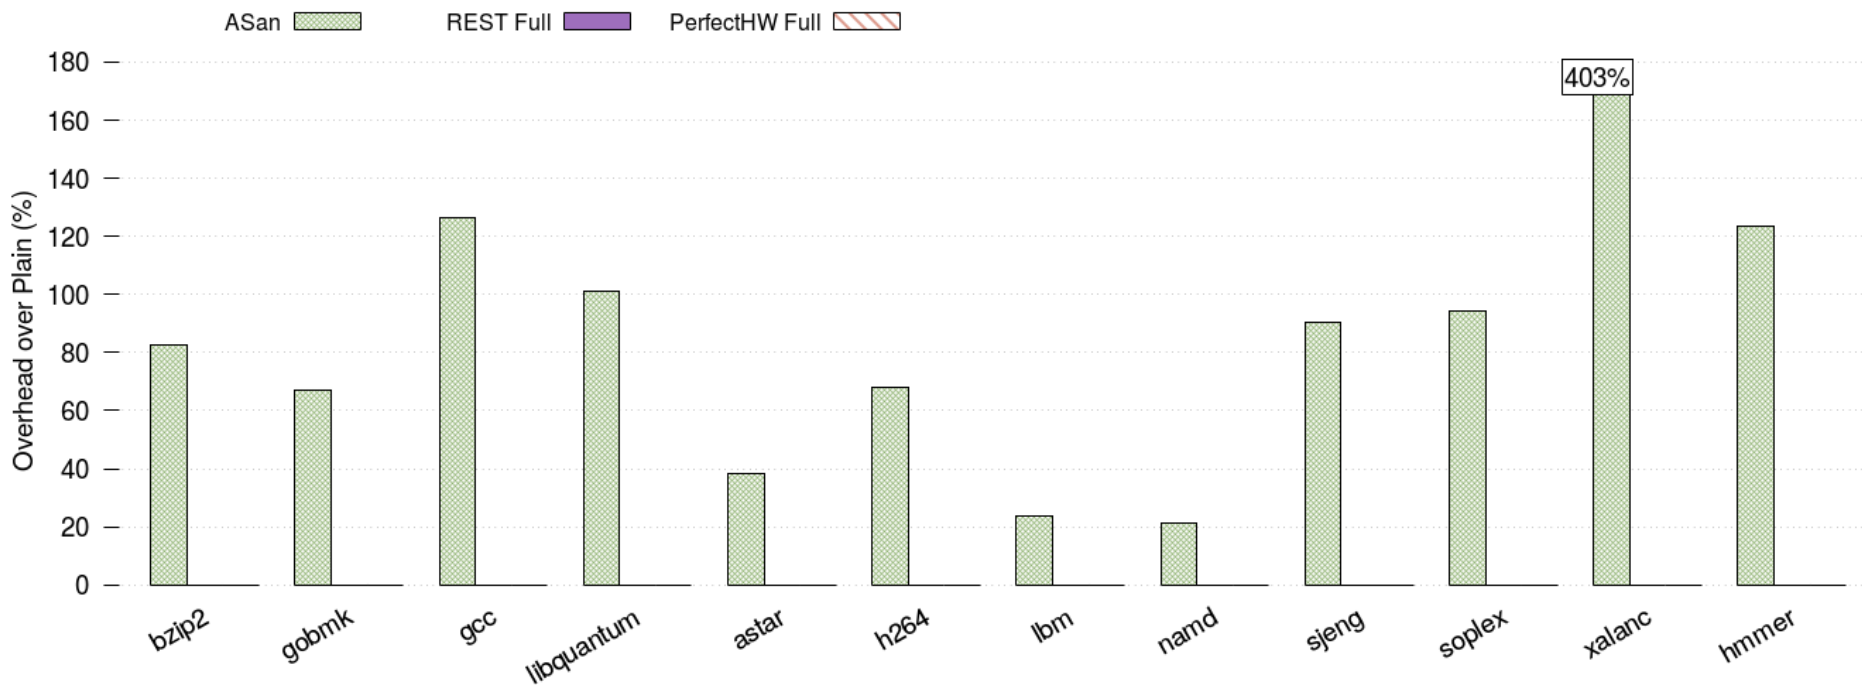
Only update token bit

Split regular match logic

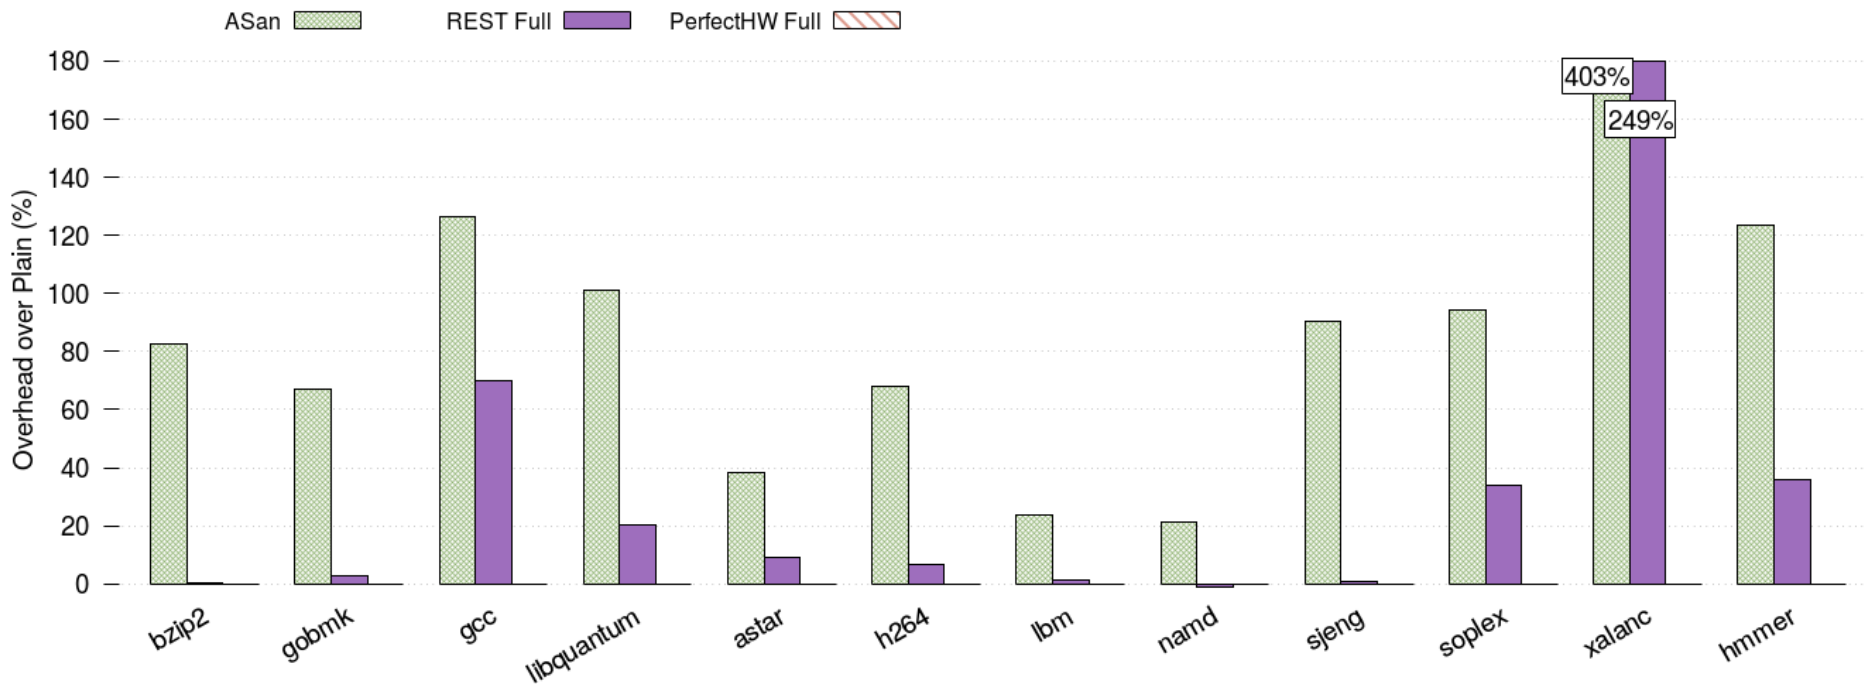


REST Overhead

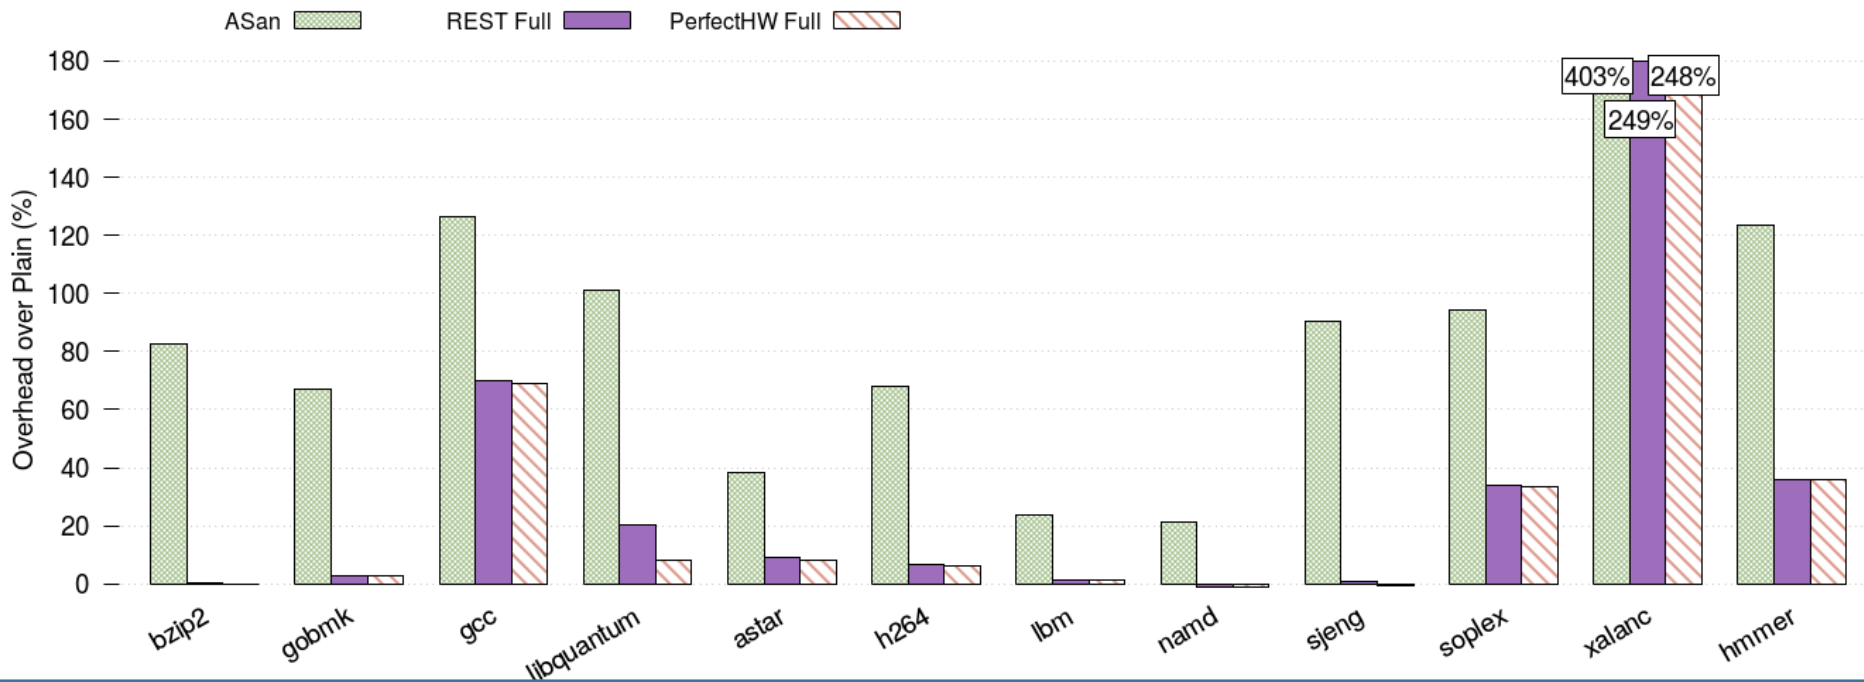
REST Performance



REST Performance



REST Performance



REST primitive overhead near-zero.
Software overhead mostly from **allocator**.

To conclude...

REST: Hardware/software mechanism to detect common memory safety errors

- Low overhead, low complexity hardware implementation
- Heap safety for legacy binaries

22-90% faster than comparable software solution on SPEC CPU

Questions?